

# Data Structures and Processing

**Dr. Carlos Henrique Brandt**  
[cbrandt@constructor.university](mailto:cbrandt@constructor.university)

# Image Processing

# References

- “Image Processing with Python”, The Data Carpentries,  
<https://datacarpentry.org/image-processing>



# Introduction

Images are an important source of information, humans have always relied on photography to registry life.

- Image processing is used for
  - detecting astronomical objects
  - detecting diseases
  - estimating populations
  - driving cars

# Introduction

- Morphometrics: counting and measuring objects (sizes, shapes)



# Image Basics

*How are images represented in digital formats?*


- Pixels and arrays
  - bit and bytes
  - Monochromatic and RGB
- Image formats
  - BMP, JPEG, TIFF
- Compression
  - lossy and lossless
- Metadata

# Image Basics

## *Pixels and arrays*

Images are stored as rectangular arrays of discrete “picture elements”, otherwise known as *pixels*.

Each pixel can be thought of as a single (square) *point of colour*.




# Image Basics

*Pixels and arrays*

Here is an image of an “8” in 15 pixels (5,3):

This image is represented by array:

```
[[0. 0. 0.]  
 [0. 1. 0.]  
 [0. 0. 0.]  
 [0. 1. 0.]  
 [0. 0. 0.]]
```





# Image Basics


*Pixels and arrays*

Here is the “8” transformed into “0”:

```
image[2,1] = 1
```

This image is represented by array:

```
[[0. 0. 0.]  
 [0. 1. 0.]  
 [0. 1. 0.]  
 [0. 1. 0.]  
 [0. 0. 0.]]
```




# Image Basics

## *Pixels and arrays*

There are a couple of aspects to notice from these simple examples:

- The colours associated to “0” and “1” were arbitrary.
  - Visualization tools will often decide for a colour palette if none is explicitly defined.
- The (axes) coordinates system start from the top-left corner.
  - X-coordinates increase to the right
  - Y-coordinates increase downwards




# Image Basics

## *Pixels and arrays*

There are a couple of aspects to notice from these simple examples:


- The colours associated to “0” and “1” were arbitrary.
  - Visualization tools will often decide for a colour palette if none is explicitly defined.
- The (axes) **coordinates system start from the top-left corner.**
  - X-coordinates increase to the right
  - Y-coordinates increase downwards



# Image Basics

## Colours

The tool we're using to visualize these images is Matplotlib, which use the Viridis colour palette by default, and normalize the values in the array to span the colour range.




# Image Basics

## *Colours*

Let's see what Matplotlib does if we add more non-zero values to our (5,3) array, and stretch the values from 0 to 255:

```
> image *= 128  
> image[2, :] = 255
```

```
[[0.  0.  0. ]  
 [0. 128. 0. ]  
 [255. 255. 255.]  
 [0. 128. 0. ]  
 [0.  0.  0. ]]
```




# Image Basics

## *Colours*

...the same image, through a grayscale colormap:

```
> plt.imshow(image, cmap=plt.cm.gray)
```

```
[[0.  0.  0. ]  
 [0. 128. 0. ]  
 [255. 255. 255.]  
 [0. 128. 0. ]  
 [0.  0.  0. ]]
```



# Image Basics


## *Colours and arrays*

Images represented in 2-dimensional are called *monochromatic* images.

The colours associated to those values (by Matplotlib, for instance) are fake colours given by a colormap. There are many palettes one can use, see:

- [matplotlib.org/gallery/color/colormap\\_reference](https://matplotlib.org/gallery/color/colormap_reference)

True colour images are represented by 3-dimensional arrays, where each colour channel is represented by a 2D array.




# Image Basics

## *Colours and arrays*

The most common standard for colour images is the additive RGB where each pixel is represented by three values: Red, Green, Blue

The following (Numpy) array will create the image we see here:

```
> rgen = np.random.RandomState(2021)
> img = rgen.randint(0,255,size=(4,4,3))
> plt.imshow(img)
```



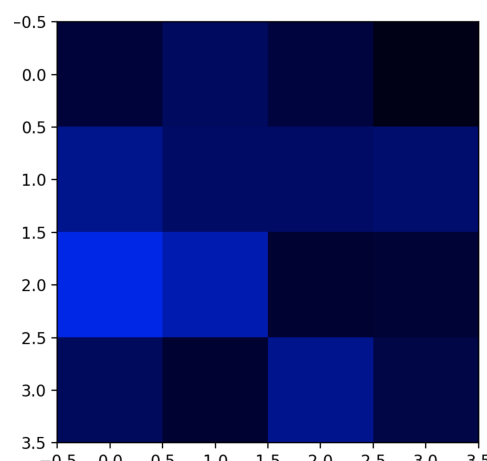





# Image Basics

## *Colours and arrays*

Each of those colours are the result of combining the values of each (R,G,B) channel:

```
> red_channel = img * [1, 0, 0]
> green_channel = img * [0, 1, 0]
> blue_channel = img * [0, 0, 1]
```



# Image Basics

*Colours, bits and bytes*

RGB images are also called 24-bit RGB images, each channel is represented by an 8-bit – or 1 byte – number (from 0-255).

A byte – 8 bits – can store 256 integer values. The amount of bits used to represent an image is called the depth.

















A monochromatic (integer) image is an 8-bit image, meaning the pixels can store 256 values.

Pixels in a (24-bit) RGB images can represent 16 million colours.

# Image Basics

*Colours, bits and bytes*

RGB colour table:

Color name	RGB triplet	Color
Red	(255, 0, 0)	
Lime	(0, 255, 0)	
Blue	(0, 0, 255)	
White	(255, 255, 255)	
Black	(0, 0, 0)	
Gray	(128, 128, 128)	
Fuchsia	(255, 0, 255)	
Yellow	(255, 255, 0)	
Aqua	(0, 255, 255)	
Silver	(192, 192, 192)	
Maroon	(128, 0, 0)	
Olive	(128, 128, 0)	
Green	(0, 128, 0)	
Teal	(0, 128, 128)	
Navy	(0, 0, 128)	
Purple	(128, 0, 128)	

# Image Basics


## *Image formats*

- **BMP**: *Device-Independent Bitmap* are non-compressed file formats capable of storing 8-bit, 16-bit, or 24-bit (colour depth) images.
- **JPEG**: *Joint Photographic Experts Group* file is a compressed lossy format capable of storing 24-bit RGB images. The amount of compression when storing an image can be adjusted.
- **PNG**: *Portable Network Graphics* file supports lossless compression and can store 8-bit (gray), 24-bit (RGB), and 32-bit (RGBA) images.
- **TIFF**: *Tag Image File Format* can be uncompressed, or compressed using either lossy or lossless. TIFs can stored more than one image.


# Image Basics

## *Image formats*

Effects of compression in an image:




Uncompressed TIF



# Image Basics

## *Image formats*

Effects of compression in an image:



Compressed JPG



# Image Basics

## *Image metadata*

Some image formats (like JPEG and TIFF) support the inclusion of metadata (along with the image data).

Metadata is textual information about the image itself, such as time when the image was created, location, technical information about the camera taking the picture, etc.

# Working with Numpy

- Image I/O
  - read image
  - write image
- Matplotlib
  - show image
- Numpy
  - image as array



# Working with Numpy

## *Image I/O*

To read and write images we will use the *imageio* Python library, and visualize it using Matplotlib.

```
> import imageio.v3 as iio
> import matplotlib.pyplot as plt
>
> filename = "maize-root-cluster.jpg"
> image = iio.imread(uri=filename)
> plt.imshow(image)
>
> # save as a TIF file
> iio.imwrite("roots.tif", image)
```



# Working with Numpy

## *Manipulating colours*

Once we have the image in memory (as an array) we can manipulate the pixels. For example, let's convert all (dark) background into black.

```
> import imageio.v3 as iio
> import matplotlib.pyplot as plt
>
> filename = "maize-root-cluster.jpg"
> image = iio.imread(uri=filename)
>
> image[image < 128] = 0
>
> plt.imshow(image)
```




# Working with Numpy

## *Histograms*

Intensity/Colour histograms play an important role in image processing, giving a deeper understanding of image through their light distribution.

```
> filename = "plant-seedling.jpg"
> image = iio.imread(uri=filename,
                        mode="L")
>
> counts, edges = np.histogram(
    image.flatten(),
    bins=256,
    range=(0,256))
>
> plt.plot(edges[:-1], counts)
```




# Working with ~~Numpy~~ Matplotlib

## *Histograms*

... we can also use Matplotlib (directly) to plot our image's histogram:

```
> filename = "plant-seedling.jpg"
> image = iio.imread(uri=filename,
                        mode="L")

>
> _ = plt.hist(
    image.flatten(),
    bins=256,
    range=(0,256)
)
```

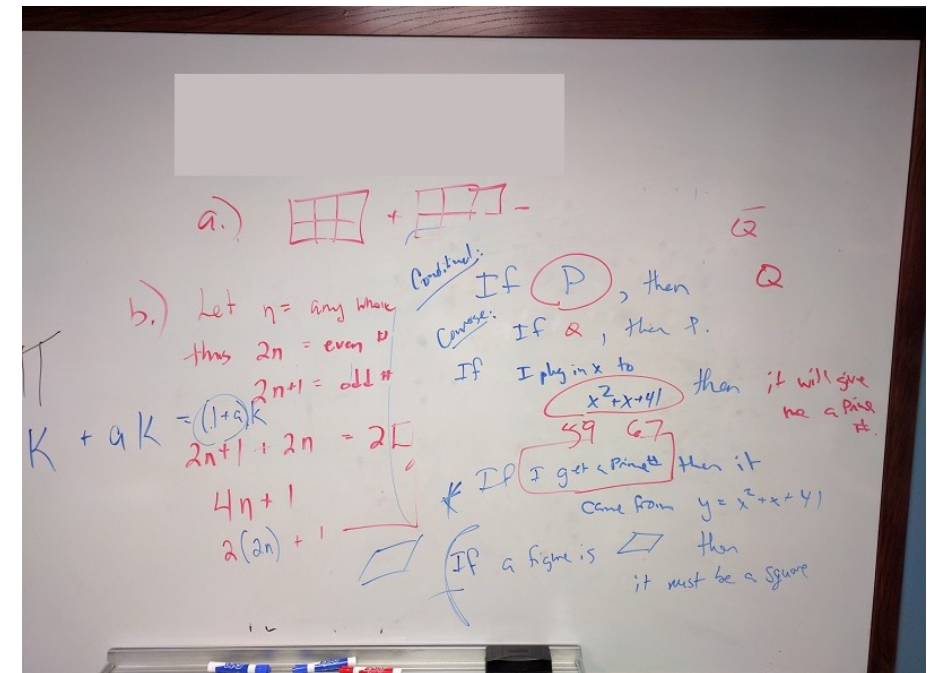
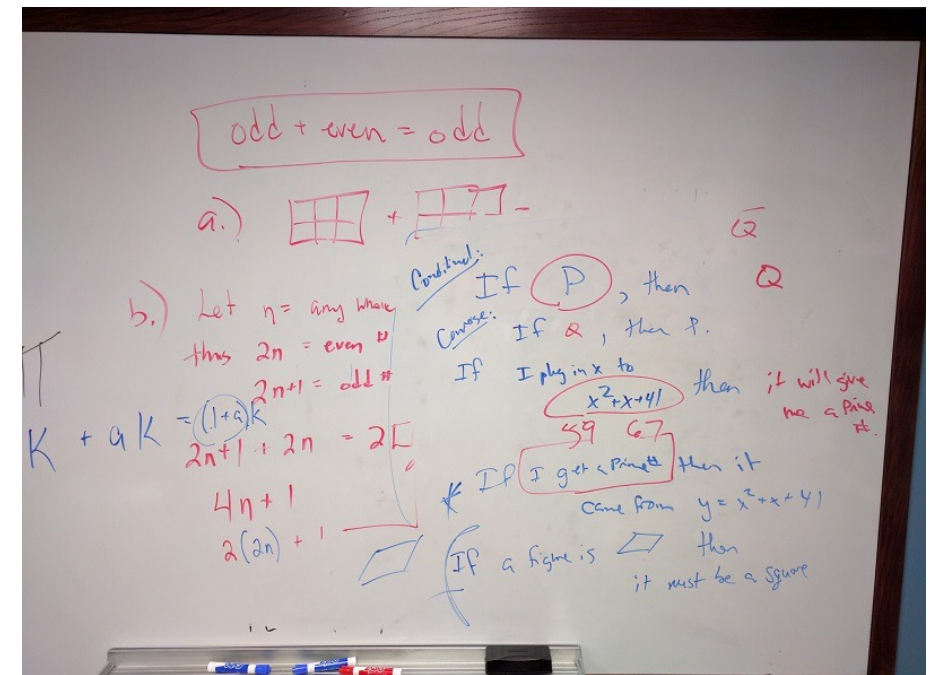


# Working with Numpy

## Manipulating colours

We can also clip parts of an image:

```
> filename = "board.jpg"
> image = io.imread(uri=filename)
> fig, ax = plt.subplots()
> ax.imshow(image)
>
> clip = image[60:151, 135:481, :]
> color = image[330, 90]
>
> image[60:151, 135:481] = color
> fig, ax = plt.subplots()
> ax.imshow(image)
```



# Working with Skimage

## *Manipulating colours*

We can use another library, scikit-image, to do image transformation. For example, to convert a colour image into grayscale.

```
> from skimage.color import rgb2gray
>
> filename = "maize-root-cluster.jpg"
> image = io.imread(uri=filename)
>
> gray_img = rgb2gray(image)
>
> plt.imshow(gray_img, cmap="gray")
```



# Working with Skimage

## *Resizing images*

Skimage provides tools for image transformations such as resizing and conversion between datatypes (eg, from float to unsigned integer).

```
> from skimage.transform import resize
> from skimage.util import img_as_ubyte
>
> new_shape = (
    image.shape[0] // 2,
    image.shape[1] // 2,
    image.shape[2]
)
> small = resize(image, new_shape)
> small = img_as_ubyte(small)
```





# Working with Skimage

## *Masking*

In many situations in image processing we want to retain or manipulate only some parts of the image. We do that by using *masks*.

Masks are Boolean arrays where the “pixels” in the regions (or patches) of interest are valued True (False) whereas the rest are False (True).

The way the pixels (or array elements) of interest are defined varies according to the problem, you can, for instance, define the elements through colour/value selection (like we did when we turned all dark pixels into black “0”), or you can use a pre-defined list of [y,x] positions.



# Working with Skimage

## *Masking*

If we want to mask a certain region of the image and we want to create a list of the  $[y,x]$  elements in that area we can use skimage drawing tools to accomplish that.

Suppose we want to retain only the roots of the seeds in the following image:

```
> filename = "maize-seedlings.tif"
> image = iio.imread(uri=filename)
>
> fig, ax = plt.subplots()
> ax.imshow(image)
```



# Working with Skimage

## *Masking*

Suppose we want to retain only the roots of the seeds in the following image, we can define the rectangle covering the region through a mask and manipulate the image as follows:

```
> from skimage.draw import rectangle
>
> mask = np.ones(image.shape[0:2],
                  dtype=bool)
>
> rr, cc = rectangle(start=(357, 44),
                     end=(740, 720))
>
> mask[rr,cc] = False
> image[mask] = 0
```



# Segmentation

Segmentation is the process of dividing an image in regions of interest, it is central to identification of objects/content in images. The whole segmentation process is composed by the following stages:

- pre-processing: prepare the image (eg, smoothing, normalization);
- processing: apply strategy to split image into regions or layers;
- post-processing: filter and identify objects of interest.

# Segmentation

## *Blurring images*

The blurring of images is usually a necessary step as it smooths differences (eg, object border) between neighbour pixels, and distribute imperfections (eg, noise) among nearby pixels.


In signal processing, “blurring” is called a *low-pass filter*. On the other hand, *high-pass filters* are used to *enhance* differences (eg, borders) between nearby pixels.

When we talk about “blurring” an image, we usually mean the use of a *Gaussian filter*. A Gaussian filter is very similar (effect) to a simple *mean filter*, but the gaussian preserves better the image information.


# Segmentation

## *Gaussian blur*


*To refresh our memory, how Gaussian curves/filter look like:*



Curves with varied mean/std-dev values



3-dimensional surface




7x7 array kernel

# Segmentation

*Gaussian blur*

Animations of blurring kernel being applied:





# Segmentation

## *Gaussian blur - Skimage*

Skimage has a whole module dedicated to filters, *gaussian* is one. See how results change according to different *sigma* (ie, filter size) values.

```
> from skimage.filters import gaussian  
>  
> image = iio.imread("gaussian.png")  
> blur = gaussian(image, sigma=3)
```







# Segmentation

## *Gaussian blur – 3d View*

Just so we have a better view of the effect of blurring,  
a 3D view of the “petri dish” image (after a blur):





# Segmentation

## *Thresholding*

The simplest – yet powerful – way to split an image's content is through the use of threshold method.



In this method, a value “ $t$ ” in image's intensity range is used to mark the pixels (with intensity) “above  $t$ ” and the pixels “below  $t$ ”.

Such value “ $t$ ” may be chosen through different methods and is usually niche specific. In most of the cases/images, though, the method known as *Otsu's method* is usually a good first-guess.

# Segmentation

## *Thresholding*

For example, consider the histogram section of this presentation, we used the “seedling” image to draw the following histogram:





# Segmentation

## *Thresholding*

Let's say that every pixel with intensity below  $t=200$  should be turned black and every pixel above, white (in reality, False/True)



```
> image = iio.imread("seedling.jpg",  
                      mode="L")  
  
> t = 200  
> mask = image > t
```



# Segmentation

## *Otsu's thresholding*

In some cases, may not be that simple to define the threshold value “t”.




# Segmentation

## *Otsu's thresholding*

The [Otsu's method](#) automate the process of defining “t” by looking for the (intensity) value that maximizes the information on each class.

```
> from skimage.filters import (  
    threshold_otsu  
)  
  
>  
> image = iio.imread("tree.jpg",  
    mode="L")  
  
>  
> t = threshold_otsu(image)  
>  
> mask = image > t  
> plt.imshow(mask)
```




# Segmentation

## *Otsu's thresholding*

The [Otsu's method](#) automate the process of defining “t” by looking for the (intensity) value that maximizes the information on each class.

```
> from skimage.filters import (  
    threshold_otsu  
)  
  
>  
> image = iio.imread("tree.jpg",  
    mode="L")  
  
>  
> t = threshold_otsu(image)  
>  
> mask = image > t  
> plt.imshow(mask)
```







# Segmentation

## *Otsu's thresholding*




The [Otsu's method](#) automate the process of defining “t” by looking for the (intensity) value that maximizes the information on each class.



# Segmentation

*Object identification*

We learned separate foreground (objects) from background pixels:






# Segmentation

## *Object identification*

We learned separate foreground (objects) from background pixels.

The question now is how to identify each one ('A', 'B', 'C', ...).

To do that we use a technique called Connected Component Analysis.




# Segmentation

## *Object identification*

Skimage provides identification of objects that we can use to *label* a binary image into individual objects.

```
> from skimage.measure import label
> from skimage.color import label2rgb
>
> lbl_image = label(binary_mask,
                    connectivity=2)
>
> rgb_lbl_image = label2rgb(lbl_image)
>
> plt.imshow(rgb_lbl_image)
```



# Segmentation

## *Morphometrics*

Finally, we get to measure each of the identified objects. To do so, we use the image return from the labelling function (`lbl_image`) so we proceed to analyse and filter and select the objects we want.

```
> from skimage.measure import regionprops
>
> obj_features = regionprops(lbl_image)
>
> obj_lbl_area = [(o["label"],o["area"]) for o in obj_features]
>
> print(obj_lbl_area)
[(1, 317791.0), (2, 15.0), (3, 1.0), (4, 5.0), (5, 521173.0), (6, 494815.0), (7,
514227.0), (8, 6.0), (9, 137.0), (10, 254455.0), (11, 16.0), (12, 1.0), (13, 1.0),
(14, 57.0), (15, 338003.0), (16, 264905.0), (17, 8.0), (18, 8.0)]
```

# Homework

# Homework

- The data files used in the following exercises are provided at this week's course material as `data.zip` (Moodle/Teams), or directly from Figshare:
  - [https://figshare.com/articles/dataset/Data\\_Carpentry\\_Image\\_Processing\\_Data\\_beta\\_/19260677](https://figshare.com/articles/dataset/Data_Carpentry_Image_Processing_Data_beta_/19260677)
- The following (Python) libraries are necessary:
  - Scikit-Image
  - Matplotlib
  - Numpy

See <https://datacarpentry.org/image-processing/setup/> for details.

# Homework

## *Basics*

1. Load “8” image, file “data/eight.tif”, and simply change the value of pixels so you have what looks like a 5 instead of an 8.  
Display the image and print out the matrix (ie, array) as well.
2. Suppose we represent colours as triplets (R,G,B) of integer values between [0:255], what are the following colours:
  - a. (255, 0, 0)
  - b. (0, 255, 0)
  - c. (0, 0, 255)
  - d. (255, 255, 255)
  - e. (0, 0, 0)
  - f. (128, 128, 128)
3. How many colours can be represented in a 24-bit RGB system?

# Homework

## *Working with Numpy*

1. Turn the (white) background of image “data/sudoku.png” into gray.
2. Consider image “data/maize-root-cluster.jpg”, clip the roots/plant from the (larger) image, display it, and save it as “data/clip.tif”.
3. Make the colour histogram (ie, counts for R, G, and B, in the same) of image “data/tree.jpg”.

# Homework

## *Working with Numpy*

1. Turn the (white) background of image “data/sudoku.png” into gray.
2. Consider image “data/maize-root-cluster.jpg”, clip the roots/plant from the (larger) image, display it, and save it as “data/clip.tif”.
3. Make the colour histogram (ie, counts for R, G, and B, in the same) of image “data/tree.jpg”.



# Homework

## *Blurring and Thresholding*


1. Try the Gaussian filter function from Skimage. See the (visual) effects of different combinations of 'sigma', 'truncate', 'multichannel' parameters using "gaussian-original.png" image.
2. Create a black/white image (aka, *binary image*) using thresholding method to separate the objects (*True*) from the background (*False*) in image "data/shapes-02.jpg".
3. Create a black/white image (aka, *mask*) to separate (foreground) bacteria colonies from the background in image "colonies-1.tif":
  - a. Using a manually defined (through histogram) "t" value;
  - b. Using an automatically defined "t" value through Otsu's method.

# Homework

## *Connected Component Analysis*

1. How many objects/regions would be defined in the following image/array, using

- a. 4-neighborhood (1-jump) strategy?
- b. 8-neighborhood (2-jumps) strategy?



2. Re-apply the segmentation and labelling steps we saw in the “Segmentation/Object Identification” slides but, now, *before transforming the labelled image into RGB* (‘label2rgb’), re-label the objects in (a copy of) the image using their respective areas.
3. Also, does the number of objects detected match your expectations? Why, and how to fix it?

# Homework

- Solve the problem – Capstone Challenge – at:
  - <https://datacarpentry.org/image-processing/09-challenges/index.html>

You will apply all the steps we went through in this lesson to count the number of bacteria colonies in petri-dish images

# References

- “Image Processing with Python”, The Data Carpentries,  
<https://datacarpentry.org/image-processing>